

PARALLEL AND GPU PROGRAMMING IN PYTHON

PRACE Training Course

Ben Czaja, Sagar Dolas, Mohsen Safari
HPC advisors, SURF




SURF

© Annemiek van de

Outline

- Using GPUs in Python
- CUDA Programming Model
- CUDA Execution Model
- Examples:
 - Vector (1D array) Addition
 - Matrix (2D array) Addition
 - Matrix Multiplication
- Optimization Tips

Accessing to GPUs in Python



PyCUDA gives you easy, Pythonic access to NVIDIA's CUDA parallel computation API.

<https://documen.tician.de/pycuda/>

CUDA Programming Model

- Introduced by NVIDIA in 2006, Compute Unified Device Architecture
- General purpose programming model that leverages the parallel compute engine in NVIDIA GPUs
- An extension of C language
- CUDA programs are CPU-GPU programs:
 - CPU part is called *host*
 - GPU part is called *kernel*

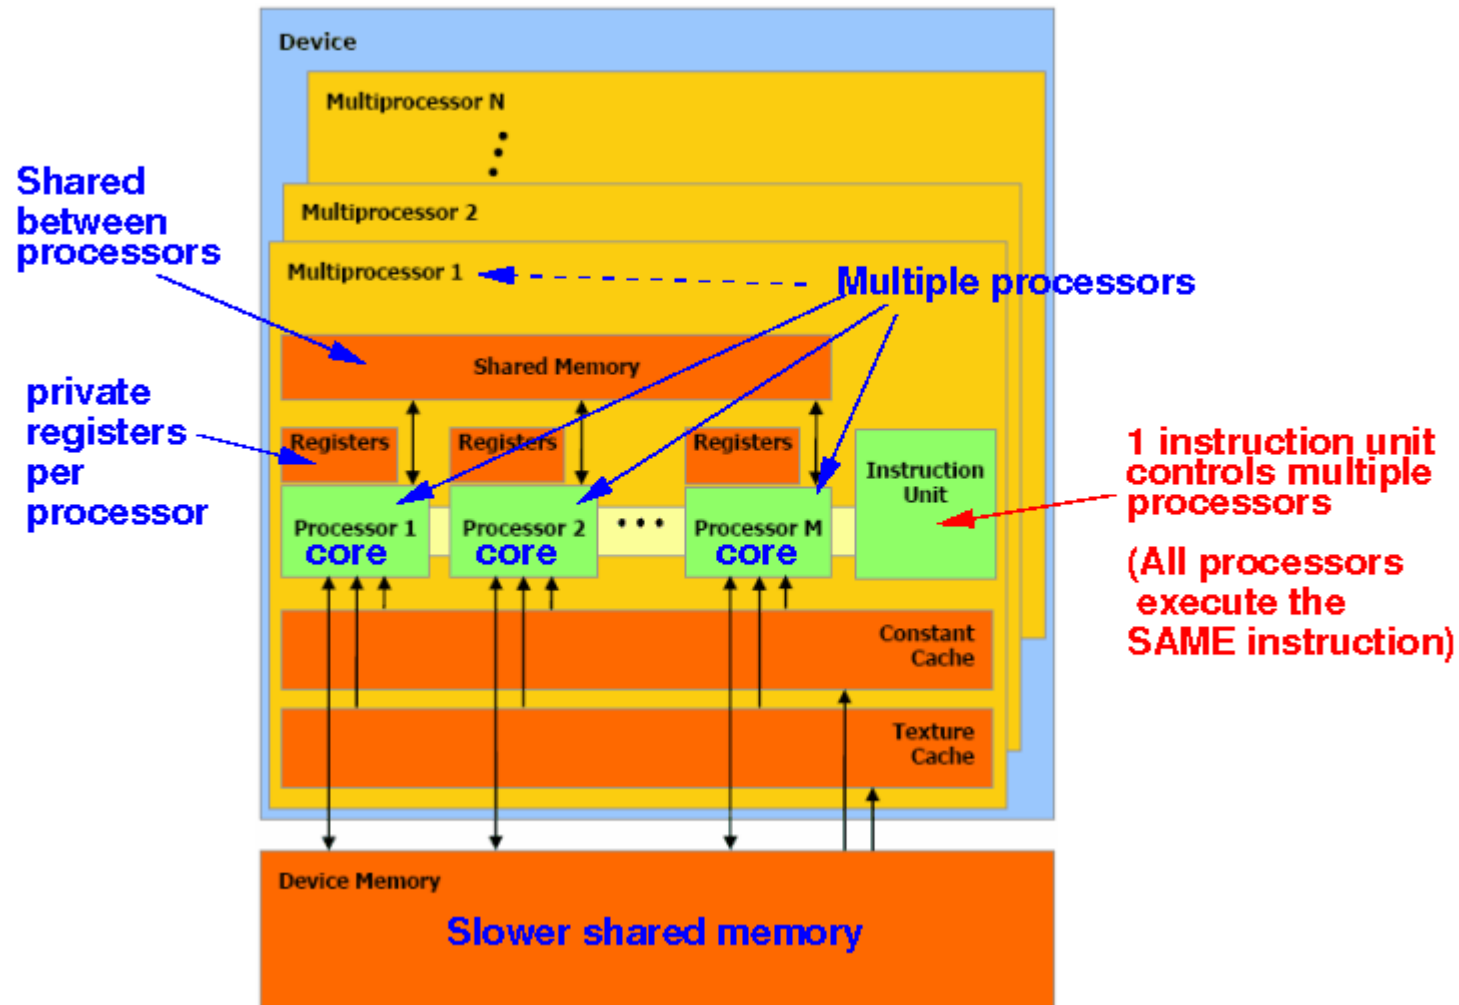
CUDA Programming Model

To execute any CUDA program, there are three main steps:

- Copy the input data from host memory to device memory, also known as host-to-device transfer
- Call the kernel from host and execute the GPU program
- Copy the results from device memory to host memory, also called device-to-host transfer

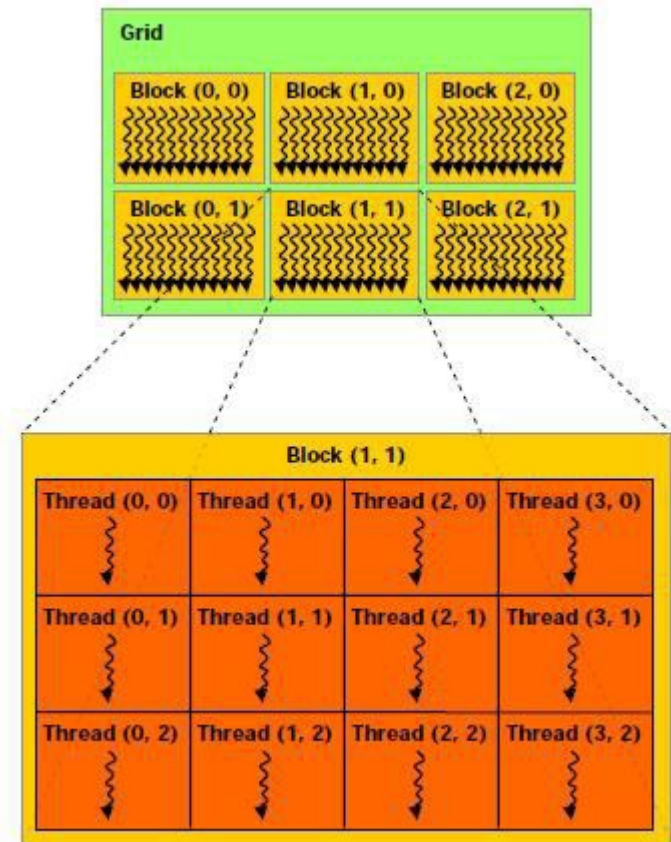
NVIDIA GPU Hardware

GPU device:



CUDA Programming Model

- Threads are organized into two hierarchical levels:
 - Threads are grouped into *blocks*
 - Blocks are grouped into *grids*
- Blocks and grids can be 1D, 2D and 3D



CUDA Programming Model

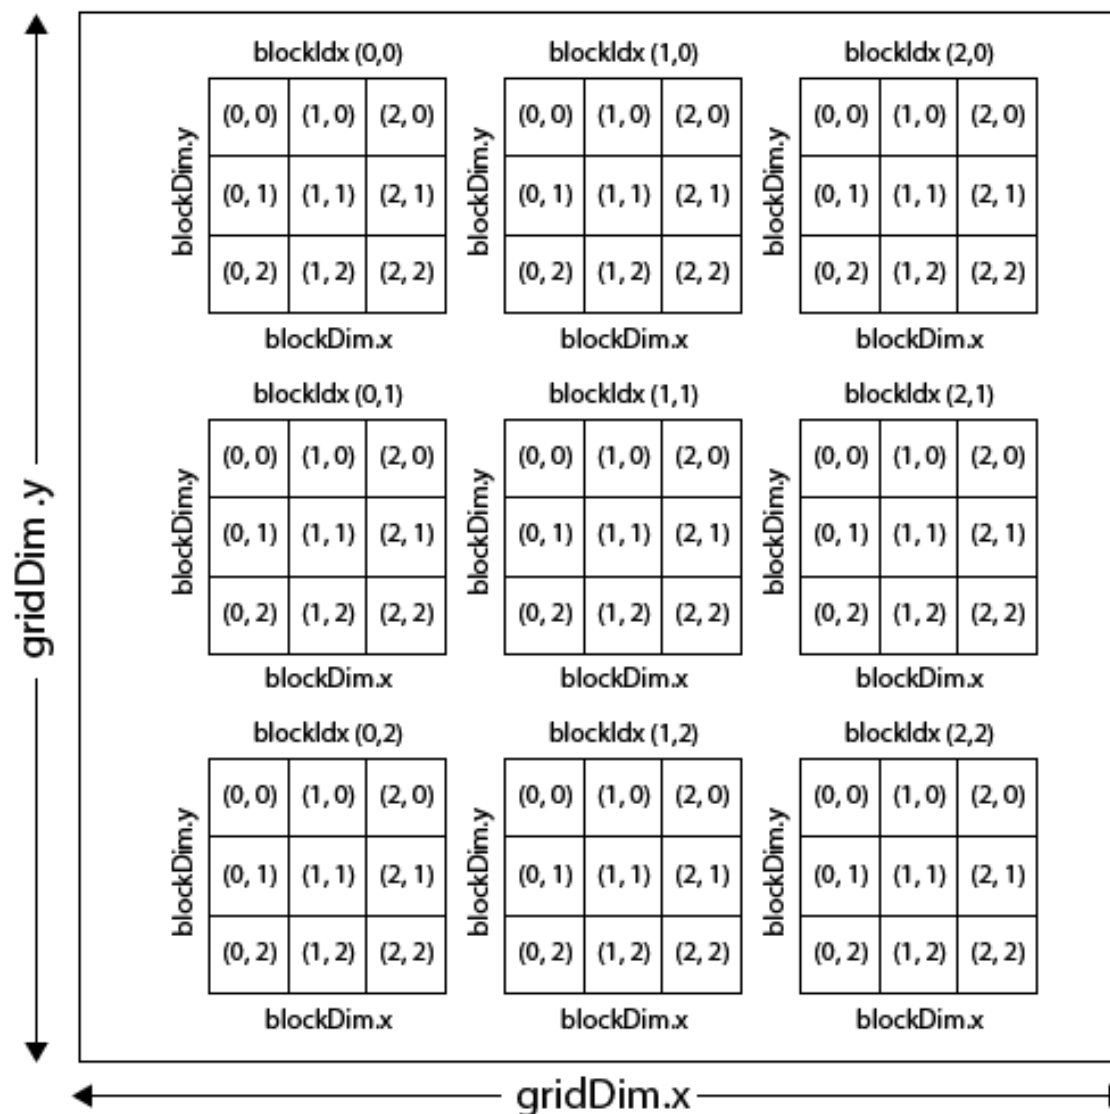
■ Built-in functions:

- Dimension:
 - `gridDim.x`, `gridDim.y`, `gridDim.z`
 - `blockDim.x`, `blockDim.y`, `blockDim.z`
- Index:
 - `blockIdx.x`, `blockIdx.y`, `blockIdx.z`
 - `threadIdx.x`, `threadIdx.y`, `threadIdx.z`

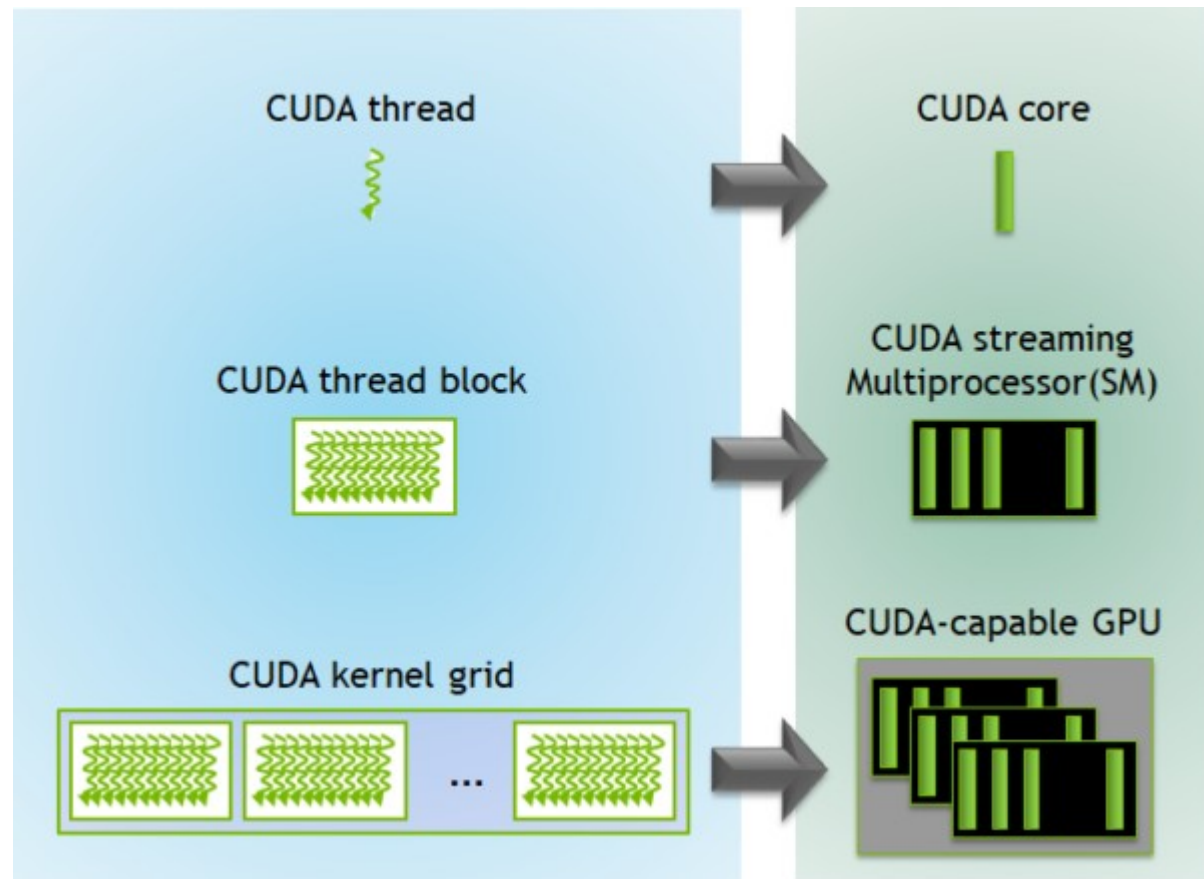
CUDA Programming Model

CUDA Grid


- gridDim.x = 3
- gridDim.y = 3
- blockDim.x = 3
- blockDim.y = 3



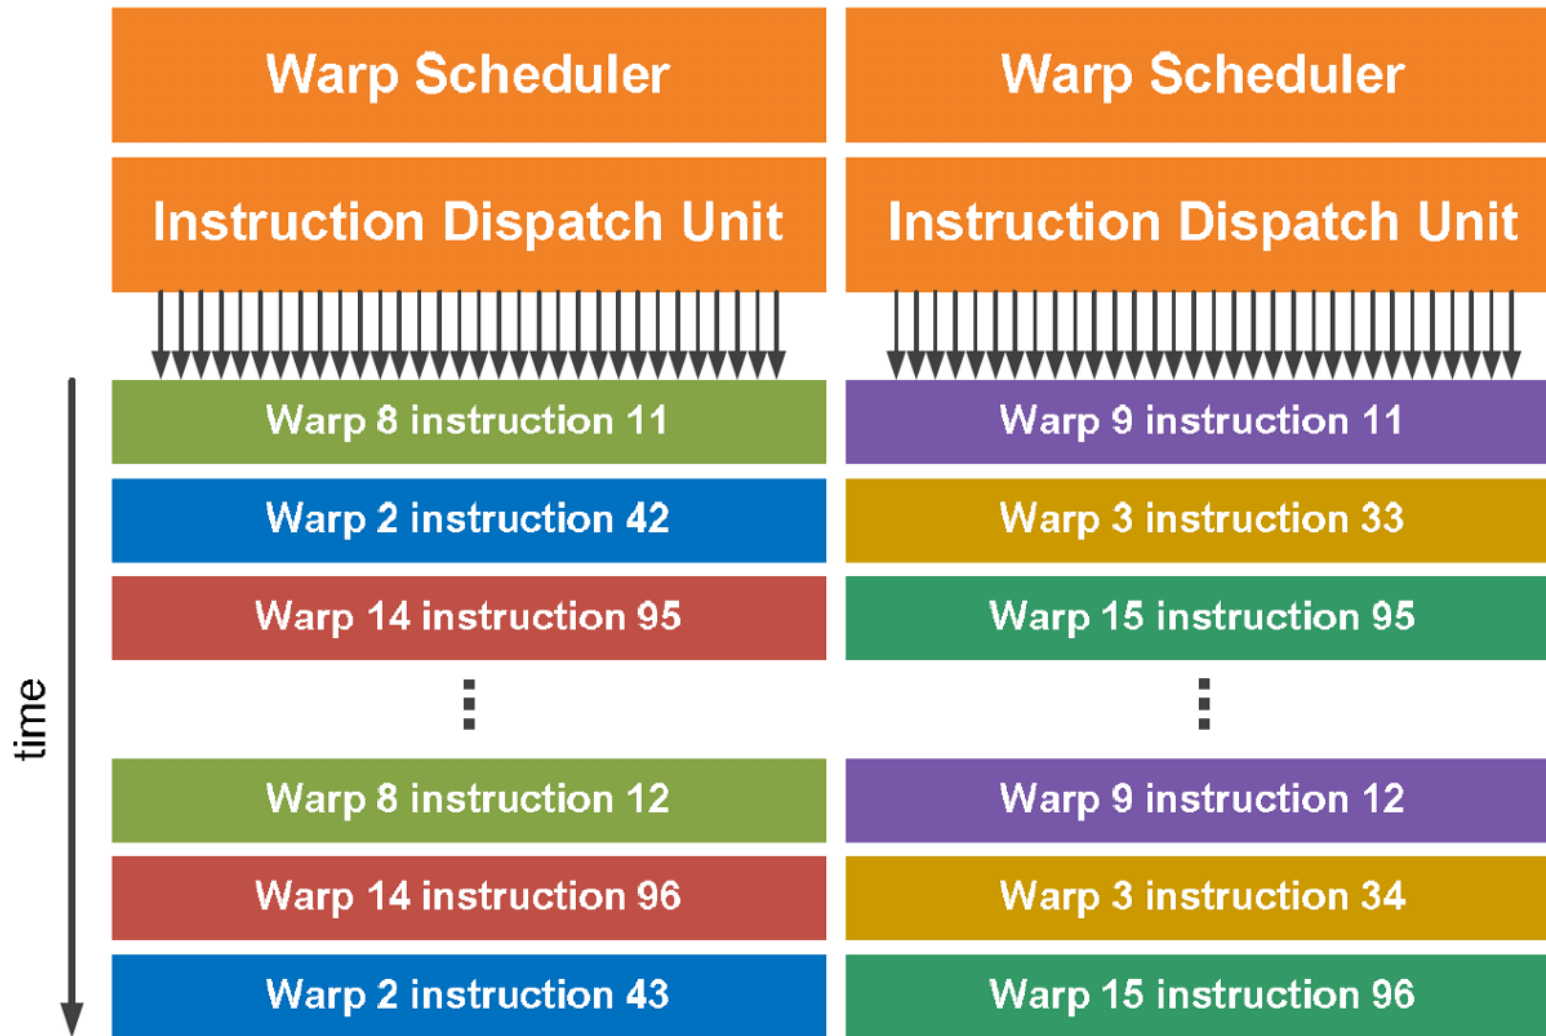
CUDA Execution Model



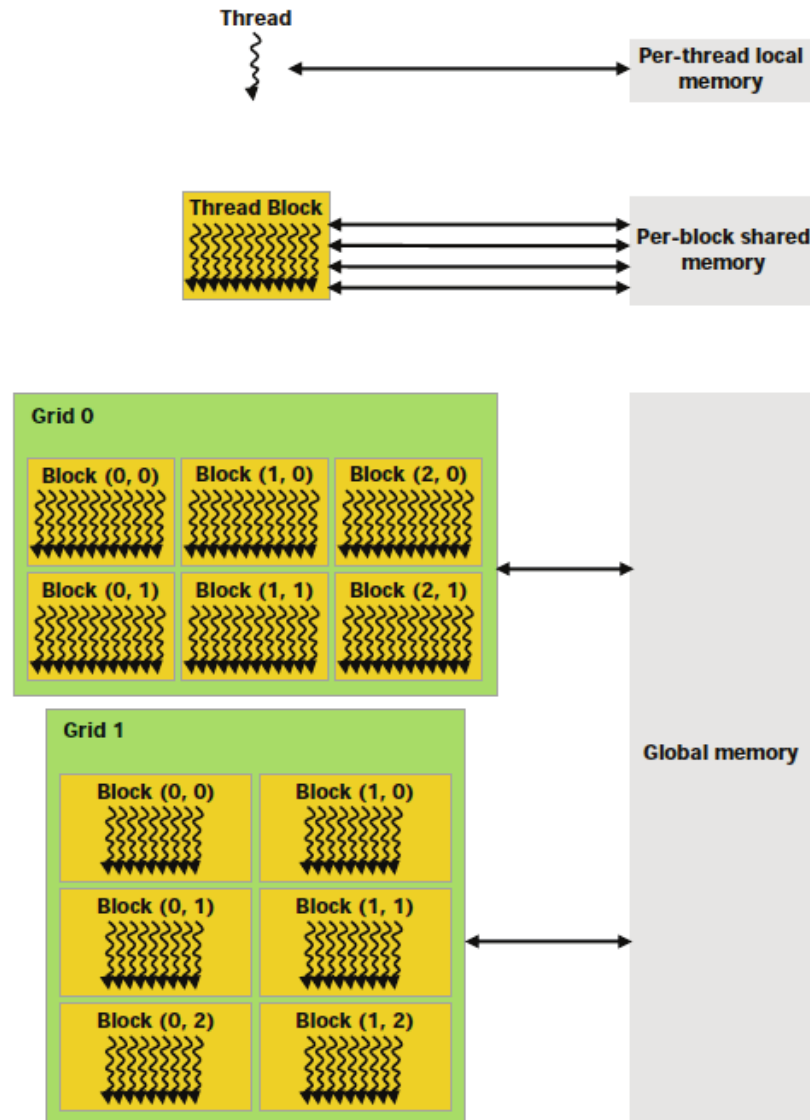
CUDA Execution Model



CUDA Execution Model



CUDA Execution Model



Synchronization in CUDA

- There is a mechanism to synchronize all threads in a block:
 - Built-in function `__syncthreads()`
- There is no mechanism to synchronize all threads across all blocks

GPU Node

- 4 NVIDIA Titan RTX GPUs per node
 - Multiprocessors: 72
 - Streaming cores: 4608
 - Global memory: 24 GB
- One node is shared among 16 people
- One GPU is shared among 4 people
- Note that you have around 5 GB memory:
 - Matrix $(35,000 * 35,000) = 35,000 * 35,000 * 4 \approx 5 \text{ GB}$
 -

First Example:

Parallel Vector (1D array) Addition in PyCUDA

Calculate Global Index (1D grid, 1D block)

Global Thread ID

0	1	2	3	4	5	6	7	8	9	10	11	12	13	14	15	16	17	18	19	20	21	22	23	24	25	26	27	28	29	30	31
threadIdx.x								threadIdx.x								threadIdx.x								threadIdx.x							
0	1	2	3	4	5	6	7	0	1	2	3	4	5	6	7	0	1	2	3	4	5	6	7	0	1	2	3	4	5	6	7
blockIdx.x = 0								blockIdx.x = 1								blockIdx.x = 2								blockIdx.x = 3							

- Global Thread ID: $\text{blockIdx.x} * \text{blockDim.x} + \text{threadIdx.x}$
- For global thread ID 26:
 - $\text{blockIdx.x} = 3$
 - $\text{blockDim.x} = 8$
 - $\text{threadIdx.x} = 2$
 - Global thread ID = $3 * 8 + 2 = 26$

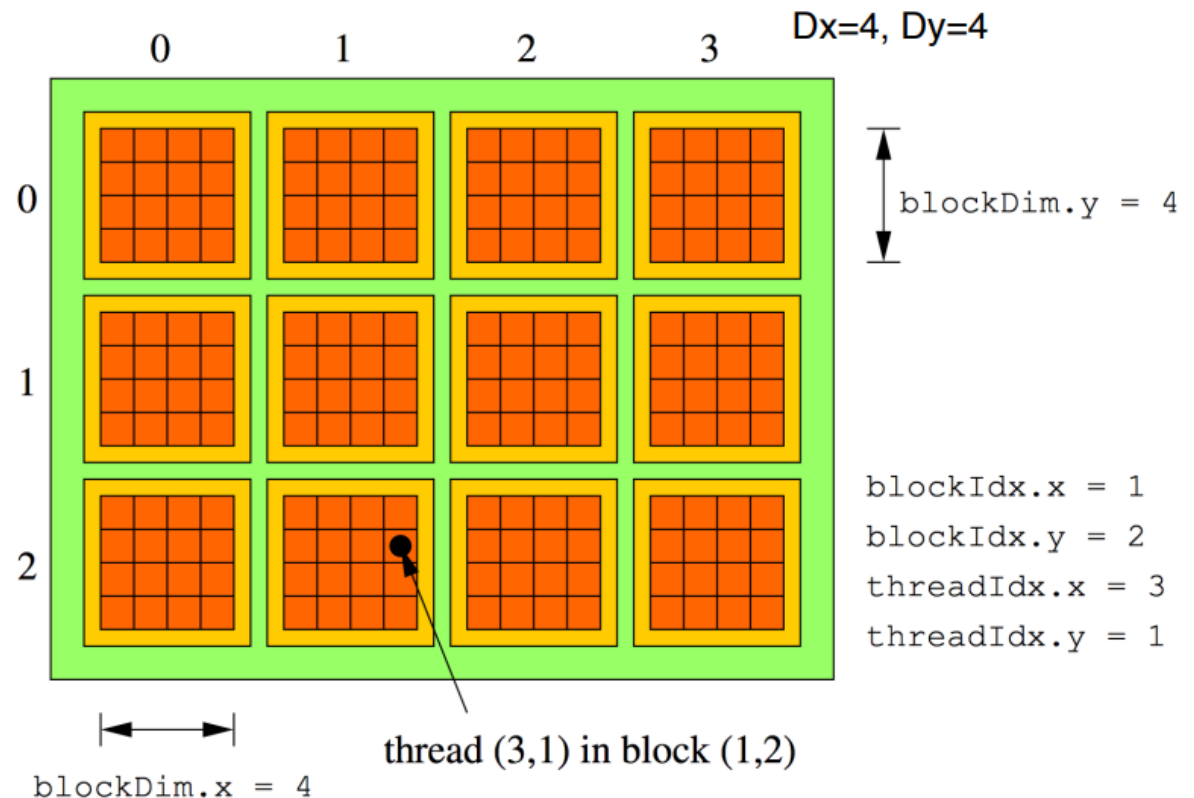
Automatic Data Transfer

- Automatic data transfer using PyCUDA driver:
 - `In()`
 - `Out()`
 - `InOut()`
- PyCUDA programs become simpler

Second Example:

Parallel Matrix (2D array) Addition in PyCUDA

Calculate Global Index (2D grid, 2D block)



Matrix 12*16

Global Thread ID:

- $\text{row} = \text{blockIdx.y} * \text{blockDim.y} + \text{threadIdx.y} = 2 * 4 + 1 = 9$
- $\text{column} = \text{blockIdx.x} * \text{blockDim.x} + \text{threadIdx.x} = 1 * 4 + 3 = 7$

Row-Major Flattening of a Matrix

- Matrix 3*3
- For each element (row, col):
 - New ID = row * (No of col) + col
- For instance element “5” in location (1, 2):
 - New ID = 1 * 3 + 2 = 5

How we see a 2D array

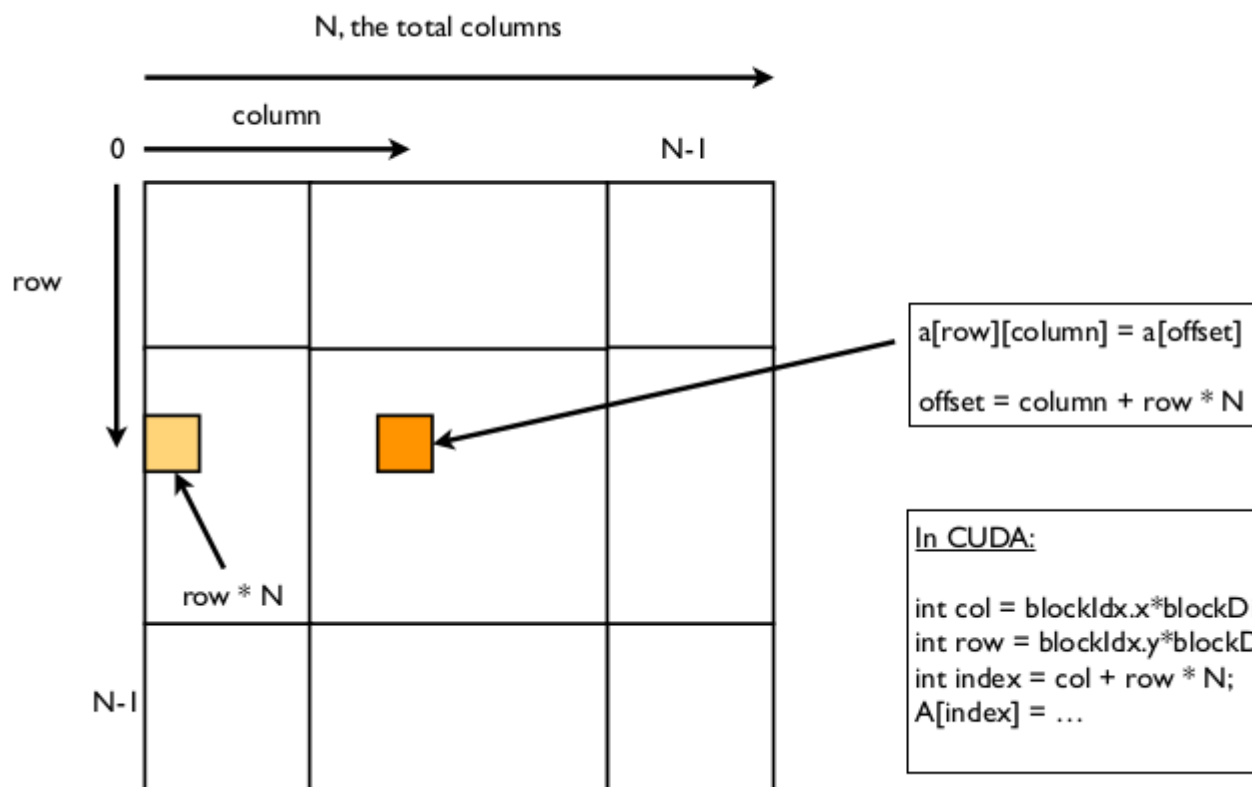
0	1	2
3	4	5
6	7	8



How it's stored in memory

0	1	2	3	4	5	6	7	8
---	---	---	---	---	---	---	---	---

Row-Major Flattening of a Matrix



In CUDA:

```
int col = blockIdx.x*blockDim.x+threadIdx.x;  
int row = blockIdx.y*blockDim.y+threadIdx.y;  
int index = col + row * N;  
A[index] = ...
```


Third Example:

Parallel Matrix (2D array) Multiplication in
PyCUDA

Sequential Matrix Multiplication

```
/* ijk */
for (i=0; i<n; i++) {
  for (j=0; j<n; j++) {
    sum = 0.0;
    for (k=0; k<n; k++)
      sum += a[i][k] * b[k][j];
    c[i][j] = sum;
  }
}
```

Inner loop:




$$\begin{bmatrix} a & b \\ c & d \end{bmatrix}_A \times \begin{bmatrix} e & f \\ g & h \end{bmatrix}_B = \begin{bmatrix} ae + bg & af + bh \\ ce + dg & cf + dh \end{bmatrix}_C$$

Parallel Matrix Multiplication

```
int k, sum = 0;

int col = threadIdx.x + blockDim.x * blockIdx.x;
int row = threadIdx.y + blockDim.y * blockIdx.y;

if(col < width && row < width) {
    for (k = 0; k < width; k++)
        sum += a[row * width + k] * b[k * width + col];
    c[row * width + col] = sum;
}
```




Time comparison

Time comparison between

- PyCUDA
- Numpy.matmul()
- @ operator

Further Optimization

- Matrix A: $M \times W$
- Matrix B: $W \times N$
- Matrix C: $M \times N$
- Assume $W = 32$
- Assume $\text{blockDim.x} = 32$
- Assume $\text{blockDim.y} = 32$



Exercise 1

```
__shared__ int a_shared[32][32];  
__shared__ int b_shared[32][32];  
a_shared[threadIdx.y][threadIdx.x] = A[...];  
b_shared[threadIdx.y][threadIdx.x] = B[...];  
__syncthreads();  
For (int k = 0; k < 32; k++)  
    temp += a_shared[threadIdx.y][k] * b_shared[k][threadIdx.x];  
C[id] = temp;
```

Exercise 2

- How to extend it to any size W ?
- How is the performance now?


Optimization

There are different ways to optimize CUDA codes:


- Number of threads per block
- Workload per thread
- Total work per thread block
- Correct memory access and data locality
- ...

Tips for Optimization

- Global Memory Access:



- Coalesced




- Non-coalesced

Tips for Optimization

- Avoid Warp Divergence:

```
...  
if ( threadIdx.x < 16 )  
{  
    ... A ...  
}  
else  
{  
    ... B ...  
}  
...
```



Tips for Optimization

■ Use shared memory in two cases:

- When threads in a block need to share data
- When there are repeated accesses to one location in global memory
 - In this case, it is possible to use register as local memory to each thread

Questions

Thank you for participating! Any questions?